


## Chapter 24

### Single-Source Shortest Paths

Let  $G$  be a weighted directed graph where the weight attached to each edge is the cost to traverse the edge.

A path from vertex  $u$  to vertex  $v$  is a sequence of one or more edges,  $\{(v_1, v_2), (v_2, v_3), \dots, (v_{r-1}, v_r)\}$ , in  $G$  where  $v_1 = u$  and  $v_r = v$ . The cost of the path is the sum of the weights of the edges in the sequence.

The shortest path weight from  $u$  to  $v$  is the minimum cost of all paths from  $u$  to  $v$ . If there is no path from  $u$  to  $v$  then the shortest path weight is  $\infty$ .



Shortest paths are not necessarily unique, and neither are shortest-path trees. The above figure shows a weighted, directed graph and two shortest-path trees with the same root.

Four different shortest path problems can be posed on the weighted directed graph,  $G$ .

1. From a given source vertex,  $s$ , find the shortest path weights to all other vertices in  $G$ .
2. To a given destination vertex,  $t$ , find the shortest path weights from all other vertices in  $G$ .
3. Given two vertices,  $u$  and  $v$ , find the shortest path from  $u$  to  $v$ .
4. For every pair of vertices,  $u$  and  $v$ , find the shortest path weight from  $u$  to  $v$ .

**Negative weights:** Suppose  $G$  has one or more edges with negative weights. Such an edge actually reduces the cost of any path using it. Suppose there's a cycle using edges whose total cost is negative. A path can run around the cycle many, many times and get any negative cost desired. If the cycle can be reached from the source vertex,  $s$ , and if the destination vertex,  $d$ , can be reached from the cycle then the shortest-path weight from  $s$  to  $d$  is  $-\infty$ .

INITIALIZE-SINGLE-SOURCE( $G, s$ )

```

1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = NIL$ 
4   $s.d = 0$ 



```

RELAX( $u, v, w$ )

```

1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 

```



## The Bellman-Ford algorithm


- Allows negative-weight edges
- Computes  $v.d$  and  $v.\pi$  for all  $v \in V$
- Returns TRUE if no negative-weight cycles reachable from  $s$ , FALSE otherwise

BELLMAN-FORD( $G, w, s$ )


```

1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE


```




(a)




(b)



(c)



(d)



(e)

The nested for loops relax all edges  $|V| - 1$  times. The Bellman-Ford algorithm runs in time  $\Theta(|V||E|)$ , since the initialization in line 1 takes  $\Theta(|V|)$  time, each of the  $|V| - 1$  passes over the edges in lines 2-4 takes  $\Theta(|E|)$  time, and the for loop of lines 5-7 takes  $\Theta(|E|)$  time.

## Dijkstra's algorithm

Dijkstra's algorithm solves shortest-path problem 1 for directed weighted graph with non-negative edge weights. It is a greedy algorithm and similar to MST-Prim. Starting at the source vertex,  $s$ , it grows a tree,  $T$ , that eventually spans all vertices reachable from  $s$ . If  $v$  is a vertex in  $G.V$  then  $v.d$  holds the shortest-path weight from  $s$  to  $v$ . During the algorithm there are three classes of vertices: T-vertices that are in  $T$ ; A-vertices not in  $T$  but adjacent to T-vertices; and X-vertices that are not adjacent to T-vertices.

Vertices are added to  $T$  in distance order; first  $s$ , then the vertex closest to  $s$ , then the next closest, etc. If  $u$  is a T-vertex then  $u.\pi$  holds the tree-parent of  $u$  and  $u.d$  holds the shortest path weight from  $s$  to  $u$ .

If an A-vertex,  $v$ , is adjacent to one T-vertex,  $u$ , then  $v.\pi = u$  and  $v.d = u.d + w(u, v)$ . If an A-vertex is adjacent to multiple T-vertices,  $u_1, u_2, \dots, u_r$ , then

$$v.d = \min(u_i.d + w(u_i, v)) \text{ for } 1 \leq i \leq r$$

to select the T-vertex,  $u_k$ , that gives the shortest  $s$ -to- $v$  path and  $v.\pi = u_k$ .

If  $v$  is an X-vertex then  $v.\pi = NIL$  and  $v.d = \infty$ .

When a vertex,  $u$ , becomes a T-vertex,  $Adj[u]$  is scanned for A-vertices and X-vertices. If A-vertex,  $v$ , is adjacent to  $u$  then  $v.d$  is compared to  $u.d + w(u, v)$  to see if  $v.d$  and  $v.\pi$  need to be updated. If X-vertex,  $v$ , is adjacent to  $u$  then  $v$  is changed to an A-vertex,  $v.\pi = u$ , and  $v.d = u.d + w(u, v)$ .


A-vertices and X-vertices are stored in a priority queue,  $Q$ , ordered by their  $d$ -variables with maximum priority given to the vertex with least distance,  $d$ . As with MST-PRIM, EXTRACT-MIN( $Q$ ) extracts the vertex with least distance.

DIJKSTRA( $G, w, s$ )


```

1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}(Q)$ 
6       $S = S \cup \{u\}$ 
7      for each vertex  $v \in G.Adj[u]$ 
8          RELAX( $u, v, w$ )


```




(a)




(b)




(c)



(d)



(e)



(f)

Like MST-PRIM, DIJKSTRA runs in  $O(|E| \lg |V|)$  time.